

A decorative graphic on the left side of the slide, consisting of a network of light blue lines and circles, resembling a circuit board or a neural network diagram. The lines are of varying thickness and the circles are of varying sizes, creating a complex, branching pattern that extends from the top to the bottom of the frame.

# DEEP REINFORCEMENT LEARNING

RL FOR THE MODERN AGE....

The background is a dark teal gradient. In the corners, there are decorative white line art elements resembling circuit boards or neural network connections. These include straight lines, right-angle turns, and small circles at the end of the lines.

# BACKGROUND INFORMATION....

# WHY DEEP RL?


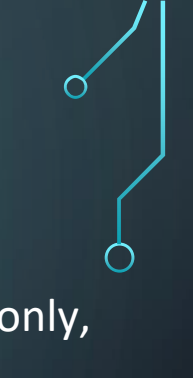
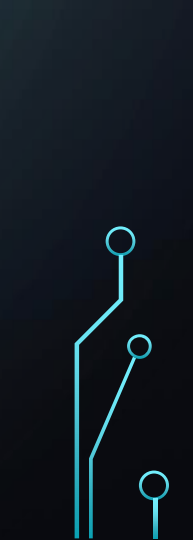
- You can frame many problems in terms of a discrete set of States, however almost all “real-world” data is continuous.
- We therefore need to use some continuous function in order to learn from continuous data
- Neural Network are learnable continuous function!

# PROBLEMS WITH USING NEURAL NETS....

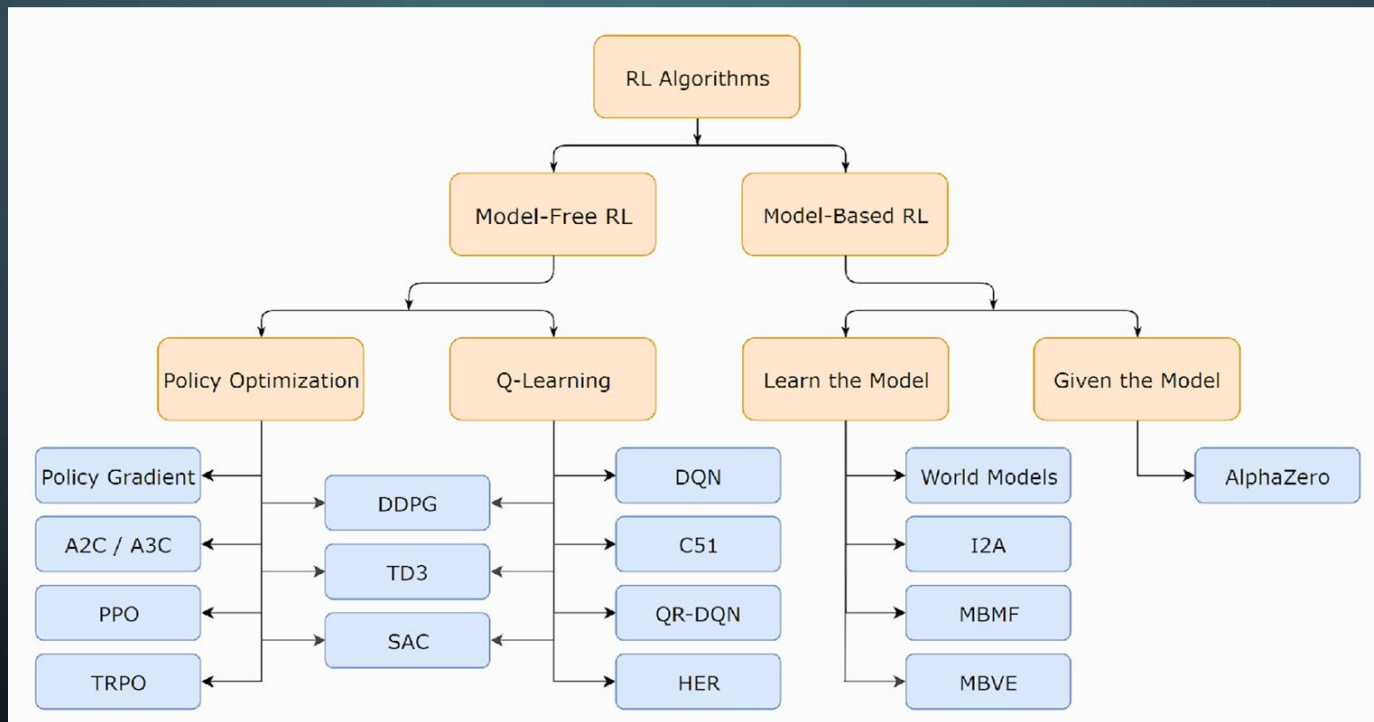
- RL the training signal is often very sparse
  - Rewards are often only given once, or far apart during a rollout
- Training signals are often dependant of each other
  - You only get some rewards after completing an initial step which gets a reward.
    - In a video game, to get to the second enemy you need to kill the first enemy first, if you never kill the first enemy you never see the second.
- The training data distribution is non-stationary
  - As the agent gets better it is likely to progress further into a task, getting itself into previously unseen States.
    - To get to the second level in a video game, you need to complete the first level!



# PROBLEMS WITH USING NEURAL NETS....

- Unlike a Look-Up table, you can't change the stored variable for a single State only, Neural Networks are parameterised equations, if you change the values of the parameters you change the outputs for ALL States!
  - Neural Networks require a LOT of data to train, and can easily over-fit to a small dataset!
- 
- 
- 

# CURRENT LANDSCAPE OF DEEP RL



# DEEP Q NETWORKS (DQN)

---

## Playing Atari with Deep Reinforcement Learning

---

**Volodymyr Mnih   Koray Kavukcuoglu   David Silver   Alex Graves   Ioannis Antonoglou**

**Daan Wierstra   Martin Riedmiller**

DeepMind Technologies

`{vlad,koray,david,alex.graves,ioannis,daan,martin.riedmiller} @ deepmind.com`

### **Abstract**



We present the first deep learning model to successfully learn control policies directly from high-dimensional sensory input using reinforcement learning. The

model is a convolutional neural network, trained with a variant of Q-learning.

# DEEP Q NETWORKS (DQN)

- One of the first instances of creating a robust and “general” Deep RL algorithm
- The “Q-Value” table is replaced with a Convolutional “Q Network”

$Q(a, s)$	$a_0$	$a_1$	$a_2$	....	$a_n$
$s_0$					
$s_1$					
$s_2$					
$s_3$					
$s_4$					
$\vdots$					
$s_m$					





# CHALLENGES WITH DEEP Q NETWORKS (DQN)

To train a “Deep Q Network” we need to turn the Q-Value update into a loss function:

$$L_i(\theta_i) = \mathbb{E}_{s,a \sim \rho(\cdot)} \left[ (y_i - Q(s, a; \theta_i))^2 \right]$$

where  $y_i = \mathbb{E}_{s' \sim \mathcal{E}} [r + \gamma \max_{a'} Q(s', a'; \theta_{i-1}) | s, a]$

We can see that target  $y_i$  is just the Q-Value update and the total loss is MSE between the current Q-Value and the update.

Here we also notice a bit of a problem! Our “target” for the Q-Value, is constructed from the output of our Q-Network! This can lead to instabilities.

# DEEP Q NETWORKS (DQN) - Solutions

- Non-Stationary/Noisy data and dependant training signals
  - Experience Replay Buffer
    - A FIFO buffer that collects and holds each “transition” experience  $(s_t, a_t, r_t, s_{t+1})$
    - ~10,000 - 1,000,000 transitions
    - At training time we perform batched gradient descent over the average loss
- Sparse Rewards/Overfitting
  - A relatively small network

# DEEP Q NETWORKS (DQN)

---

**Algorithm 1** Deep Q-learning with Experience Replay

---

Initialize replay memory  $\mathcal{D}$  to capacity  $N$

Initialize action-value function  $Q$  with random weights

**for** episode = 1,  $M$  **do**

    Initialize sequence  $s_1 = \{x_1\}$  and preprocessed sequenced  $\phi_1 = \phi(s_1)$

**for**  $t = 1, T$  **do**

        With probability  $\epsilon$  select a random action  $a_t$

        otherwise select  $a_t = \max_a Q^*(\phi(s_t), a; \theta)$

        Execute action  $a_t$  in emulator and observe reward  $r_t$  and image  $x_{t+1}$

        Set  $s_{t+1} = s_t, a_t, x_{t+1}$  and preprocess  $\phi_{t+1} = \phi(s_{t+1})$

        Store transition  $(\phi_t, a_t, r_t, \phi_{t+1})$  in  $\mathcal{D}$

        Sample random minibatch of transitions  $(\phi_j, a_j, r_j, \phi_{j+1})$  from  $\mathcal{D}$

        Set  $y_j = \begin{cases} r_j & \text{for terminal } \phi_{j+1} \\ r_j + \gamma \max_{a'} Q(\phi_{j+1}, a'; \theta) & \text{for non-terminal } \phi_{j+1} \end{cases}$

        Perform a gradient descent step on  $(y_j - Q(\phi_j, a_j; \theta))^2$  according to equation 3

**end for**

**end for**

---

## Atari 2600 video games, free online game play in your browser.



The **Atari 2600** is a video game console released in September 1977 by Atari Inc. The 2600 was typically bundled with two joystick controllers, a conjoined pair of paddle controllers, and a cartridge game — initially **Combat** and later **Pac-Man**. The Atari 2600 was wildly successful during the early 1980s. It was superseded by the **Atari 5200** and **7800** game consoles.

Our 2600 game emulator is computer, mobile phone friendly and Iphone compatible. Most games include the game manual and instructions to help you play.

(Click on Game Title to Play Game in Your Web Browser)

Sort games by:



3D Tic Tac Toe



Acid Drop



Action Force



Adventure



Air Raiders



Airlock



Alien



Alien's Return



Alpha Beam With Ernie




Amidar



Angriff Der Luftklotzen




Artillery Duel



# CODE EXAMPLE

CARTPOLE



## Deep Reinforcement Learning with Double Q-learning

Hado van Hasselt and Arthur Guez and David Silver  
Google DeepMind

### Abstract

The popular Q-learning algorithm is known to overestimate action values under certain conditions. It was not previously known whether, in practice, such overestimations are common, whether they harm performance, and whether they can generally be prevented. In this paper, we answer all these

questions. We show that overestimation is a common exploration technique (Kaelbling et al., 1996). If, however, the overestimations are not uniform and not concentrated at states about which we wish to learn more, then they might negatively affect the quality of the policy. We show that this is not the case and that Schwartz (1993) gives a simple way to avoid this problem. This leads to suboptimal policy

## Dueling Network Architectures for Deep Reinforcement Learning

Ziyu Wang  
Tom Schaul  
Matteo Hessel  
Hado van Hasselt  
Marc Lanctot  
Nando de Freitas

ZIYU@GOOGLE.COM  
SCHAUL@GOOGLE.COM  
MTTHSS@GOOGLE.COM  
HADO@GOOGLE.COM  
LANCTOT@GOOGLE.COM  
NANDODEFREITAS@GMAIL.COM

## Hindsight Experience Replay

Marcin Andrychowicz\*, Filip Wolski, Alex Ray, Jonas Schneider, Rachel Fong, Peter Welinder, Bob McGrew, Josh Tobin, Pieter Abbeel<sup>1</sup>, Wojciech Zaremba<sup>1</sup>  
OpenAI

### Abstract

One of the biggest challenges in Reinforcement Learning (RL) is the problem of sparse rewards. We present a novel technique called *Hindsight Experience Replay* (HER) that allows us to avoid the need for complicated reward engineering in an arbitrary off-policy RL algorithm and may be se

processes  
element  
is use

In spite of this, most of the approaches for RL use standard neural networks, such as convolutional networks, MLPs, LSTMs and autoencoders. The focus in these recent advances has been on designing improved control and RL algorithms, or simply on incorporating existing neural net-

## Rainbow: Combining Improvements in Deep Reinforcement Learning

Matteo Hessel  
DeepMind

Joseph Modayil  
DeepMind

Hado van Hasselt  
DeepMind

Tom Schaul  
DeepMind

Georg Ostrovski  
DeepMind

Will Dabney  
DeepMind

Dan Horgan  
DeepMind

Bilal Piot  
DeepMind

Mohammad Azar  
DeepMind

David Silver  
DeepMind

### Abstract

The deep reinforcement learning community has made several independent improvements to the DQN algorithm. However, it is unclear which of these extensions are complementary and can be fruitfully combined. This paper examines six extensions to the DQN algorithm and empirically studies their combinations. Our experiments show that the combination provides state-of-the-art performance on the Atari 2600 benchmark, both in terms of data efficiency and final performance. We also provide results from a detailed ablation study that shows the contribution of each component to overall per-



## Distributional Reinforcement Learning with Quantile Regression

Will Dabney  
DeepMind

Mark Rowland  
University of Cambridge\*

Marc G. Bellemare  
Google Brain

Rémi Munos  
DeepMind

### Abstract

In reinforcement learning an agent interacts with the environment by taking actions and observing the next state and reward. When sampled probabilistically, these state transitions, rewards, and actions can all induce randomness in the observed long-term return. Traditionally, reinforcement learn-

ing has been evaluated on the suite of benchmark Atari 2600 games (Bellemare et al. 2013).



One of the theoretical contributions of the c51 work was a proof that the distributional Bellman operator is a contraction in a maximal form of the Wasserstein metric between probability distributions. In this context, the Wasserstein

# DEEP POLICY METHODS - REINFORCE

Using some of the lessons from Deep Value-Based methods, we can start to implement some of the Policy methods as well!

Unlike DQN's we don't need a replay buffer as policy gradient methods are not as sensitive to non-stationary data (we don't learn a Q-Value).

$P(a, s)$	$a_0$	$a_1$	$a_2$	.....	$a_n$	$\sum^n P(a, s)$
$s_0$	$P(a_0, s_0)$	$P(a_1, s_0)$	$P(a_2, s_0)$		$P(a_n, s_0)$	1
$s_1$						
$s_2$						
$s_3$						
$s_4$						
$\vdots$						
$s_m$						



# DEEP POLICY METHODS - REINFORCE

We can see the REINFORCE loss function

$$L_{\theta}^{PG} = \mathbb{E}_{\pi} [\log(\pi_{\theta}(s_t, a_t)) R_t]$$

An important difference here is that we want to MAXIMIZE this value!


We want to INCREASE the log-probability if the Returns are positive etc...

Note the loss is the expectation over the current policy, so we do not use examples/trajectories from other policies (we can use a replay buffer but all the experiences need to be from the same policy- aka no update).



# CODE EXAMPLE


[CARTPOLE](#)  
[OPENAI GYM](#)



# ACTOR-CRITIC METHODS

Similar to simple policy gradient methods except we now have both an Actor (Policy) and a Critic (Value) network.

Most of the time we use a single, combined Actor/Critic Network.



# ACTOR-CRITIC POLICY GRADIENTS

The loss for the Policy is similar to before, except we now calculate the Advantage.


$$A_t = R_t - V_t$$

$$L_{\theta}^{PG} = \mathbb{E}_{\pi} [\log(\pi_{\theta}(s_t, a_t)) A_t]$$

We also need to update the Value function, we can do this with either Monte Carlo or TD updates.

# CODE EXAMPLE

CARTPOLE



# RL IN THE WILDE!

## MineRL: Towards AI in Minecraft


**tldr: Build AI in Minecraft in the **Diamond** and **BASALT** competitions!**


**tldr: Use the **MineRL** environments!**

**tldr: Get the **MineRL** dataset!**

Welcome to MineRL. We want to build agents that play Minecraft the-art Machine Learning! To do so, we have created one learning datasets with **over 60 million frames** of record. Our dataset includes a set of tasks which highlights many with current techniques: environments with lots of hierarchy are sparse, and tasks where rewards are hard to define.

Get started by installing the environment, running your first checking out tasks designed to get you started with MineRL.

**Navigate:** 

**Obtain:** 

## AlphaZero: Shedding new light on chess, shogi, and Go

90 DEC 2019

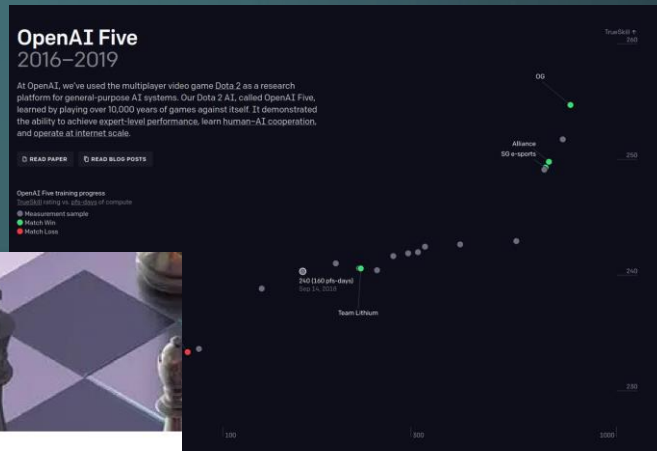
PHASE

AUTHORS


- David Silver
- Thomas Hubert
- Julian Schrittwieser
- Yann LeCun

In late 2017 we introduced AlphaZero, a single system that taught itself from scratch how to master the games of chess, shogi (Japanese chess), and Go, beating a world-champion program in each case. We were excited by the preliminary results and thrilled to see the response from members of the chess community, who saw in AlphaZero's games a ground-breaking, highly dynamic and "unconventional" style of play that differed from any chess playing engine that came before it.




Today, we are delighted to introduce the full evaluation of AlphaZero, published in the journal Science (Open Access version here), that confirms and updates those preliminary results. It describes how AlphaZero quickly learns each game to become the strongest player in history for each, despite starting its training from random play, with no in-built domain knowledge but the basic rules of the




# CAN RL DO IT ALL?


 PUBLICATIONS

SHARE



PUBLICATION LINKS


 DOWNLOAD

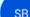
 VIEW PUBLICATION


PUBLICATION NATURE


DATE PUBLISHED  
12 MAY 2021

AUTHORS

 David Silver

 Satinder Baveja

 Doina Precup


 Richard Sutton


## Reward is Enough

### Abstract

In this paper we hypothesise that the objective of maximising reward is enough to drive behaviour that exhibits most if not all attributes of intelligence that are studied in natural and artificial intelligence, including knowledge, learning, perception, social intelligence, language and generalisation. This is in contrast to the view that specialised problem formulations are needed for each attribute of intelligence, based on other signals or objectives. The reward-is-enough hypothesis suggests that agents with powerful reinforcement learning algorithms when placed in rich environments with simple rewards could develop the kind of broad, multi-attribute intelligence that constitutes an artificial general intelligence.

### Related

 PUBLICATION  
DEEP REINFORCEMENT LEARNING  
**Deterministic Policy**

 PUBLICATION  
DEEP LEARNING  
REINFORCEMENT LEARNING  
**Mastering the game of Go**

# PROBLEMS WITH DEEP RL...

Deep RL does sound like the perfect solution for any problem! However there are some big obstacles getting in the way of mass adoption....

Deep Reinforcement Learning Doesn't Work Yet

Feb 14, 2018

June 24, 2018 note: If you want to cite an example from the post, please cite the paper which that example came from. If you want to cite the post as a whole, you can use the following BibTeX:


```
@misc{r1blogpost,
  title={Deep Reinforcement Learning Doesn't Work Yet},
  author={Irpan, Alex},
  howpublished={\url{https://www.alexirpan.com/2018/02/14/r1-hard.html}},
  year={2018}
}
```

This mostly cites papers from Berkeley, Google Brain, DeepMind, and OpenAI from the past few years, because that work is most visible to me. I'm almost certainly missing stuff from older literature and other institutions, and for that I apologize - I'm just one guy, after all.

## Introduction

Once, on Facebook, I made the following claim.

Whenever someone asks me if reinforcement learning can solve their problem, I tell them it can't. I think



[Deep Reinforcement Learning Doesn't Work Yet](#)

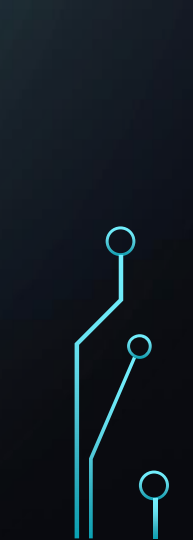
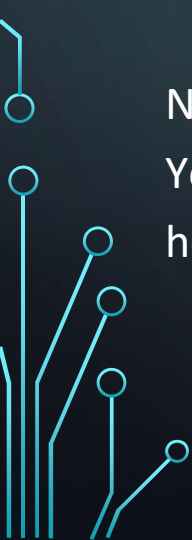


# TRAINING TIME

One of the most common problems is the time it takes to train a Deep RL agent. For example the “OpenAI Five” and Dota2 playing agent was train with over 10,000 years of game experience!!

Now imagine trying to train a robot in the real world....

You could try to train in a super fast simulation and transfer onto a real robot, however there is another problem....





# GENERALIZATION

You may have noticed one big problem with the way we train and evaluate our Deep RL agents. We train them in some environment, and then test them..... in the same environment!

A lot of recent research has shown that when testing a Deep RL agent in a slightly different environment on the same task, performance is very poor. To get a high performing and generalizable agent you not only need to provide huge amounts of data but huge amounts of variation.

Procggen

# OTHER PROBLEMS...

- RL can be very sensitive to:
  - Random initialization
  - Hyperparameters
  - Environment experiences
- Training can be unstable and agents can completely “collapse”.
  - Usually, to assess the performance of an Deep RL algorithm, many agents are trained and their performance averaged together.
- Infrastructure for training is usually more complicated than the actual RL algorithm.
  - This is why you see a lot of work in the this area by large research groups who can afford to turn a game/simulation into a training bed for RL.

# OTHER PROBLEMS...

- Hard to define a good Reward signal
  - RL will do whatever is easiest to get the reward
    - “Misaligned” agents
    - Reward hacking




# AI SAFETY?

How do we safely use Deep RL and AI in general, without knowing exactly how it will behave?

- The field of AI safety tries to come up with potential ways in which AI systems can:
  - Avoid doing what we want
  - Change its own goals
  - Become misaligned

And hopefully solve them.

[AI Safety Gridworlds](#)



THANK YOU!